PRE-DEPLOYMENT GOVERNANCE OF GENERATIVE AI

Implementing Sovereign Child-Safety Interventions in AI Coding and Conversational Agents

Strategic Project Proposal & Technical Research Paper
Prepared For: United Nations Office of Information and Communications Technology (OICT),
United Nations Office for Digital and Emerging Technologies (ODET),
United Nations Children's Fund (UNICEF), and Member States
Focus Regions: Central Europe and Spain
Date: June 2026
Status: Project Initiation, V1
Prepared By: Chandi Samlall, Nahom Ashagrea, Nakeesha Van Wyk, Nidhi Donuru

TABLE OF CONTENTS

Pre-Deployment Governance of Generative AI: Implementing Sovereign Child-Safety Interventions in AI Coding and Conversational Agents

Abstract— *The rapid proliferation of Large Language Models (LLMs) and conversational AI agents has transformed global digital interaction, introducing profound risks to vulnerable demographics, particularly children under the age of 18. This paper establishes a technical and structural framework for national governments to implement pre-deployment interventions within AI software stacks. Focusing on Central Europe and Spain as initial testing regions, the proposed model utilizes an open-source developer dashboard integrated with a child-safety decision tree. By screening foundational prompt flows against three specific categorical harms—Cyberbullying,Child Sexual Abuse and Exploitation (CSAE), and Human Trafficking—the system acts as a deterministic circuit breaker. When a threshold of harm is reached, the agent triggers an educational mitigation module that explains the context-specific limitation to the user instead of executing the prompt. This research synthesizes international digital rights frameworks, assesses structural socio-economic constraints, and maps an institutional project life cycle managed under the guidance of United Nations agencies and state regulators.*

1. Introduction

As generative artificial intelligence transitions from an emerging technology to ubiquitous computational infrastructure, its interface paradigms pose unprecedented challenges to child protection online. Traditional post-hoc algorithmic moderation often fails to neutralize threats dynamically generated by fine-tuned models or autonomous coding agents. This paper proposes a paradigm shift: pre-deployment sovereign intervention.

By establishing a standardized governance mechanism at the source-code level, national governments can collaborate directly with AI developers to embed strict, non-circumventable safety guardrails before a model is served to the public.

1.1 Objective and Scope

The primary objective of this project is the creation of an open-source source-code wrapper and corresponding evaluation dashboard. This toolkit allows developers to run automated, pre-deployment checks on conversational models built to include governance regulated by sovereign states. The initial focus is constrained geographically to Spain and the broader Central European region, utilizing the regulatory momentum of European digital mandates. The research targets three core areas of critical risk to minors:

- **Cyberbullying:** Algorithmic generation or optimization of targeted harassment, psychological abuse, or social ostracization material.
- **Child Sexual Abuse and Exploitation (CSAE):** The generation, facilitation, or solicitation of child sexual abuse material (CSAM) or grooming behaviors.

- **Human Trafficking:** Operational, logistical, or conversational patterns that facilitate the recruitment, transit, or exploitation of minors.

2. Regulatory and Ethical Frameworks

Developing an intrusive pre-deployment intervention system requires explicit alignment with existing international legal instruments. Algorithmic governance must balance proactive harm reduction with the preservation of open speech, privacy, and technical innovation.

2.1 International Human Rights and Child Guarantees

The technical architecture of the decision engine incorporates structural principles derived from the United Nations Universal Declaration of Human Rights (UDHR)[1] and the UN Convention on the Rights of the Child (UNCRC). Specifically, five articles of the UNCRC dictate the operational priorities of the algorithm's decision path:

- **Article 3 (Best Interests of the Child):** Establishes that the absolute baseline for system evaluation must prioritize minor safety over commercial model utility.
- **Articles 12 & 16 (Participation and Privacy):** Demands that the system minimize false-positive classifications to preserve a child's right to digital participation and prevent unauthorized data harvesting or surveillance during oversight.
- **Articles 19 & 34 (Protection from Harm and Exploitation):** Provides the legal mandate for deterministic state and developer intervention against cyber-violence and sexual exploitation.

2.2 Regional Compliance: Spain and Central Europe

Spain represents an optimal deployment testing ground due to its pioneering institutional frameworks, notably the Agencia Española de Supervisión de la Inteligencia Artificial (AESIA)[2]. Under AESIA guidelines and the broader European Union AI Act, providers of foundation models must adhere to rigorous risk-mitigation registries. The integration of transparent model provenance tools, such as Google Model Cards[3], ensures that safety thresholds and training configurations remain auditable to non-engineer policy stakeholders. Furthermore, ethical parameters conform to the UNESCO Recommendation on the Ethics of Artificial Intelligence[4], ensuring that algorithmic interventions respect cultural diversity and pluralism.

3. Structural and System Architecture

The technical mechanism consists of an integrated development environment (IDE) wrapper and API middleware designed to pass natural language prompts through an immutable child-safety decision pipeline before execution by the underlying LLM.

```
                [ User Prompt Input ]
                          |
                          ▼
 ┌──────────────────────────────────────────────┐
 │          Child-Safety Decision Tree          │
 │      (Cyberbullying, CSAE, Trafficking)      │
 └──────────────────────────────────────────────┘
```

```
                         │
            Does prompt trigger harm?
    ├─── YES ───► [ Trigger Mitigation Module ] ───► [ Educational
                 Interruption UI ]
└─── NO  ───► [ Underlying AI Agent Engine ] ───► [ Standard AI Output ]
```

3.1 The Child-Safety Decision Tree

The software intercepts the prompt loop at runtime. Rather than relying on black-box probabilistic classifications, the pipeline converts safety policy guidelines into an inspectable, hierarchical decision tree.

Harm Definition: A harm trigger is programmatically defined as any instance where an input prompt, inferred context, or structural system action presents a foreseeable risk of physical, psychological, economic, social, legal, privacy-related, or human-rights-related adverse consequences to an individual or group under the age of 18.

If an operational prompt falls within the probabilistic boundary of a harm vector, the model halts further inference.

3.2 Educational Interruption UI

Unlike silent failure modes or generic 'I cannot assist with that request' errors, the system triggers an educational mitigation module. This UI component renders an explicit explanation detailing precisely why the prompt path violates specific safety provisions (e.g., cross-referencing national anti-bullying or exploitation laws). This transparent, pedagogical friction deters bad actors while educating genuine users on safety boundaries.

3.3 The Developer Dashboard

To ensure auditability, the project provides a dashboard split into two main operational views:

1. **The Non-Engineer Decision Path Viewer:** Translates complex system logits and token weights into human-readable decision flows. Policy experts and legal entities can trace why a prompt was intercepted without reading raw code.

2. **Dataset-Based Evaluation Module:** Allows development teams to upload open-source validation datasets before production deployment. The system calculates performance metrics, highlighting sensitivity variations, false positives (e.g., flagging safe academic research on human trafficking), and missed risks (false negatives).

4. Empirical Performance Baseline

To parameterize the decision tree, the framework imports historical statistical medians regarding child endangerment from the United Nations Educational, Scientific and Cultural Organization (UNESCO) data repositories[5].

Harm Category	Global Context / Prevalence Baseline	Target System Sensitivity Mandate
Cyberbullying	~1 in 10 children globally impacted by digital harassment[5]; psychological bullying represents the highest risk vector within European educational sectors.	High contextual variance; must isolate linguistic peer-to-peer abuse from baseline colloquial dialogue.
Sexual Exploitation & Abuse	High prevalence globally; estimated 1 in 5 girls and 1 in 7 boys experience childhood sexual violence[6].	Zero-tolerance deterministic block; absolute priority on preventing grooming syntax.
Human Trafficking	High concentration across vulnerable, displaced, and migrant minor groups crossing regional borders.	Pattern-matching on logistical arrangements, coerced consent syntax, and identity documentation queries.

5. Challenges, Boundary Conditions, and Product Assumptions

Deploying a uniform intervention mechanism across sovereign boundaries introduces severe technical and socio-economic constraints.

5.1 System Constraints

- **Lack of Server-Side Access:** Development teams operate under the constraint of having no access to internal, proprietary server-side classifiers maintained by massive commercial AI firms. The solution must run as a modular, universal client- or edge-side layer.
- **Regulatory Fragmentations:** Divergence between country-specific local penal codes and overarching international human rights frameworks creates edge cases for what constitutes flaggable intent.
- **Socio-Economic & Cultural Complexity:** Digital syntax varies drastically across shifting socio-economic demographics. Cultural nuances among distinct ethnic groups living within identical metropolitan areas mean identical strings can carry disparate levels of psychological threat.

5.2 Product Boundary Conditions

The model assumes a normative, non-adversarial user profile. The system boundaries specifically exclude, and are not optimized to categorize, the following specialized demographics:

1. **Neurodivergent Users:** Individuals whose linguistic structure or conversational framing may inadvertently trip standard behavioral pattern-matchers.

2. **Red-Teaming Security Experts:** System testers intentionally attempting to force jailbreaks to map algorithmic limits.

3. **Academic Researchers & Policy Analysts:** Valid users inputting explicit material to analyze social trends, review legislation, or develop counter-trafficking systems.

6. Implementation Lifecycle and Stakeholder Governance

Bringing this code from conceptual framework to active national deployment requires a highly synchronized operational lifecycle split into institutional phases.

6.1 Multi-Stakeholder Coalition

- **UN Secretariat Entities:** The United Nations Office of Information and Communications Technology (OICT) and the United Nations Office for Digital and Emerging Technologies (ODET) oversee compliance with sovereign tech boundaries.
- **Open Source Oversight:** United Nations Open Source Program Offices (OSPOs) manage repository compliance, security vetting, and universal code distribution.
- **Child Advocacy Vetting:** The United Nations Children's Fund (UNICEF) provides foundational policy briefs regarding child-agent interactions[7], digital transformation strategies[8], and Digital Child Rights Impact Assessments (D-CRIA)[9] to continually refine tree logic.
- **National Deployment Actors:** United Nations Member States (specifically ministries within Spain and Central European nations), private sector tech firms, and grassroots civil societies validate localized language variations.

6.2 Phased Development Schedule

Months 1–3 (Scope Vetting): Dedicated exclusively to defining operational parameters with international stakeholders. This phase harmonizes technical boundaries with the shifting landscape of global AI governance.

Month 4 (Planning & Data Sourcing): Procurement of region-specific data pools, mapping local legislation (such as Spanish minor protection mandates), and adjusting decision algorithms to local linguistic semantics.

Month 5+ (Implementation & Testing): Deployment of the developer dashboard across selected beta cohorts in Spain and Central Europe to calibrate sensitivity and mitigate false-positive rates.

7. Conclusion

Protecting children within an AI-dominated landscape demands proactive engineering rather than retroactive moderation. By pairing a deterministic child-safety decision tree with a transparent developer dashboard, this framework provides national governments with a practical method to safeguard minor digital experiences. Orchestrated through UN agencies and applied initially within Spain and Central Europe, this architecture provides a vital equilibrium: because AI safe for kids is safe for all.

Footnotes / References

[1] United Nations. (1948). Universal Declaration of Human Rights. Retrieved from https://www.un.org/sites/un2.un.org/files/2021/03/udhr.pdf

[2] Agencia Española de Supervisión de la Inteligencia Artificial (AESIA). (2024). Guias y Regulaciones de IA. Retrieved from https://www.aesia.gob.es/ and https://aesia.digital.gob.es/en/guides

[3] Google Tools. (2020). About Model Cards. Google Research. Retrieved from https://modelcards.withgoogle.com/about

[4] UNESCO. (2021). Recommendation on the Ethics of Artificial Intelligence. Retrieved from https://www.unesco.org/en/artificial-intelligence/recommendation-ethics

[5] UNESCO. (2019). Behind the numbers: ending school violence and bullying. UNESCO Digital Library. Document ID: 366483. Retrieved from https://unesdoc.unesco.org/ark:/48223/pf0000366483

[6] UNESCO. (2026). Safe learning environments: Preventing and addressing violence in and around school. Retrieved from https://www.unesco.org/en/health-education/safe-learning-environments

[7] UNICEF Innocenti. (2026). When AI becomes a friend: Policy brief on the relationship between children and generative AI. Office of Global Insight and Policy. Retrieved from https://www.unicef.org/media/181131/file/UNICEF-When-AI-becomes-friend-policy-brief-2026.pdf

[8] UNICEF. (2022). UNICEF Digital Transformation Strategy. Retrieved from https://www.unicef.org/digitalimpact/media/1391/file/UNICEF%20DX%20Strategy_Final.pdf

[9] UNICEF. (2023). Child Rights Impact Assessment (D-CRIA) for Responsible Technology. Retrieved from https://www.unicef.org/childrightsandbusiness/workstreams/responsible-technology/D-CRIA